

▼ Anvesh Khode

PRN: 22070521021 | Practical 2: Data preprocessing and cleaning

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from scipy.stats import zscore
```

```
# ---- Load Boston Housing from CMU URL ----
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)

# Data wrangling: the dataset is split oddly across rows
X = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
y = raw_df.values[1::2, 2]

# Convert to DataFrame
feature_names = [
    "CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE",
    "DIS", "RAD", "TAX", "PTRATIO", "B", "LSTAT"
]
X = pd.DataFrame(X, columns=feature_names)
y = pd.Series(y, name="MEDV")
```

X

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT |
|-----|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|-------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 501 | 0.06263 | 0.0 | 11.93 | 0.0 | 0.573 | 6.593 | 69.1 | 2.4786 | 1.0 | 273.0 | 21.0 | 391.99 | 9.67 |
| 502 | 0.04527 | 0.0 | 11.93 | 0.0 | 0.573 | 6.120 | 76.7 | 2.2875 | 1.0 | 273.0 | 21.0 | 396.90 | 9.08 |
| 503 | 0.06076 | 0.0 | 11.93 | 0.0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1.0 | 273.0 | 21.0 | 396.90 | 5.64 |
| 504 | 0.10959 | 0.0 | 11.93 | 0.0 | 0.573 | 6.794 | 89.3 | 2.3889 | 1.0 | 273.0 | 21.0 | 393.45 | 6.48 |
| 505 | 0.04741 | 0.0 | 11.93 | 0.0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1.0 | 273.0 | 21.0 | 396.90 | 7.88 |

506 rows × 13 columns

Next steps: [Generate code with X](#) [View recommended plots](#) [New interactive sheet](#)

```
# Step 1: Handle Missing Values
# Introduce some artificial missing values (since Boston Housing usually has none)
X.iloc[0:5, 0] = np.nan
# We use **imputation (mean)** instead of deletion (loses data) or prediction (complex)
imputer = SimpleImputer(strategy="mean")
X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
```

```
# Step 2: Handle Outliers
# Outlier detection in "MEDV" target (house prices)
# We will use **removal** method (drop rows with Z-score > 3)

z_scores = np.abs(zscore(X))
X = X[(z_scores < 3).all(axis=1)]
y = y.loc[X.index]
```

```
# Step 3: Encoding categorical data
# Boston has one categorical feature: "CHAS" (Charles River dummy variable: 0 or 1)
# Since it's binary, we'll use **Label Encoding** (OneHot not needed)
# But for multi-class, OneHot would be used
# Already encoded as 0/1, so nothing more needed
```

```
# Step 4: Feature Scaling
# We'll use **Z-score normalization (StandardScaler)**
scaler = StandardScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)

print("Final cleaned and preprocessed dataset shape:", X_scaled.shape)
print()

X_scaled
```

Final cleaned and preprocessed dataset shape: (415, 13)

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LST |
|-----|-----------|-----------|-----------|------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|----------|---------|
| 0 | 0.236041 | 0.427278 | -1.254663 | 0.0 | -0.083120 | 0.503296 | -0.077906 | 0.120954 | -0.938612 | -0.607776 | -1.483794 | 0.439914 | -1.1282 |
| 1 | 0.236041 | -0.488581 | -0.553675 | 0.0 | -0.708718 | 0.251194 | 0.410713 | 0.567909 | -0.816465 | -0.943176 | -0.312853 | 0.439914 | -0.4849 |
| 2 | 0.236041 | -0.488581 | -0.553675 | 0.0 | -0.708718 | 1.501880 | -0.224135 | 0.567909 | -0.816465 | -0.943176 | -0.312853 | 0.360916 | -1.2751 |
| 3 | 0.236041 | -0.488581 | -1.273808 | 0.0 | -0.808451 | 1.195757 | -0.769819 | 1.125951 | -0.694319 | -1.067398 | 0.108686 | 0.395854 | -1.4436 |
| 4 | 0.236041 | -0.488581 | -1.273808 | 0.0 | -0.808451 | 1.439673 | -0.470228 | 1.125951 | -0.694319 | -1.067398 | 0.108686 | 0.439914 | -1.0740 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 410 | -0.498068 | -0.488581 | 0.162039 | 0.0 | 0.234212 | 0.532762 | 0.061190 | -0.700185 | -0.938612 | -0.750631 | 1.185952 | 0.344612 | -0.4029 |
| 411 | -0.501621 | -0.488581 | 0.162039 | 0.0 | 0.234212 | -0.241550 | 0.332249 | -0.797567 | -0.938612 | -0.750631 | 1.185952 | 0.439914 | -0.4942 |
| 412 | -0.498451 | -0.488581 | 0.162039 | 0.0 | 0.234212 | 1.159742 | 0.842267 | -0.858716 | -0.938612 | -0.750631 | 1.185952 | 0.439914 | -1.0261 |
| 413 | -0.488456 | -0.488581 | 0.162039 | 0.0 | 0.234212 | 0.861804 | 0.781636 | -0.745895 | -0.938612 | -0.750631 | 1.185952 | 0.372950 | -0.8962 |
| 414 | -0.501183 | -0.488581 | 0.162039 | 0.0 | 0.234212 | -0.388882 | 0.478478 | -0.686733 | -0.938612 | -0.750631 | 1.185952 | 0.439914 | -0.6797 |

415 rows × 13 columns

Next steps:

[Generate code with X_scaled](#)

[View recommended plots](#)

[New interactive sheet](#)