

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

## Data Import

```
In [4]: # Dataset: https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques

data = pd.read_csv("train.csv")
print("Dataset Loaded Successfully")
print(data.head())

Dataset Loaded Successfully
   Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0    1          60      RL       65.0     8450  Pave  NaN    Reg
1    2          20      RL       80.0     9600  Pave  NaN    Reg
2    3          60      RL       68.0    11250  Pave  NaN    IR1
3    4          70      RL       60.0     9550  Pave  NaN    IR1
4    5          60      RL       84.0    14260  Pave  NaN    IR1

   LandContour Utilities ... PoolArea PoolQC Fence MiscFeature MiscVal MoSold \
0      Lvl    AllPub   ...        0    NaN  NaN      NaN    0      2
1      Lvl    AllPub   ...        0    NaN  NaN      NaN    0      5
2      Lvl    AllPub   ...        0    NaN  NaN      NaN    0      9
3      Lvl    AllPub   ...        0    NaN  NaN      NaN    0      2
4      Lvl    AllPub   ...        0    NaN  NaN      NaN    0     12

   YrSold SaleType SaleCondition SalePrice
0    2008      WD      Normal    208500
1    2007      WD      Normal    181500
2    2008      WD      Normal    223500
3    2006      WD     Abnorml   140000
4    2008      WD      Normal    250000

[5 rows x 81 columns]
```

```
In [5]: features = ['OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF']
target = 'SalePrice'
df = data[features + [target]].dropna()
print("\nSelected Columns:")
print(df.head())

Selected Columns:
   OverallQual  GrLivArea  GarageCars  TotalBsmtSF  SalePrice
0            7       1710         2        856    208500
1            6       1262         2       1262    181500
2            7       1786         2        920    223500
3            7       1717         3        756    140000
4            8       2198         3       1145    250000
```

## Train-Test Split

```
In [6]: X = df[features]
y = df[target]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## Training the model

```
In [7]: model = LinearRegression()
model.fit(X_train, y_train)
print("\nModel Training Complete")
print("Intercept:", model.intercept_)
print("Coefficients:")
for col, coef in zip(features, model.coef_):
    print(f"{col}: {coef:.2f}")

Model Training Complete
Intercept: -94422.1245122955
Coefficients:
OverallQual: 23766.30
GrLivArea: 42.80
GarageCars: 19560.82
TotalBsmtSF: 28.40
```

```
In [8]: y_pred = model.predict(X_test)
```


```
In [9]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("\nEvaluation Metrics:")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R² Score: {r2:.3f}")

Evaluation Metrics:
```

Mean Absolute Error (MAE): 25446.05  
 Mean Squared Error (MSE): 1602914819.44  
 Root Mean Squared Error (RMSE): 40036.42  
 R<sup>2</sup> Score: 0.791

```
In [10]: residuals = y_test - y_pred
plt.figure(figsize=(4,4))
sns.histplot(residuals, kde=True, color="blue")
plt.title("Residual Distribution")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```



```
In [11]: plt.figure(figsize=(4,4))
plt.scatter(y_test, y_pred, color="green", alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted House Prices")
plt.show()
```

